

Term/Concept	Description
Top-P	0 to 1 value — cumulative probability, balances randomness and accuracy, 1 is least random.
Temperature	Controls randomness—higher values lead to diverse, creative outputs; lower values make results more deterministic.
Epochs	Number of iterations on training a model—more is generally better, but too many can lead to overfitting.
AWS Rekognition	Computer vision/image recognition, used for object detection, facial analysis, and content moderation.
AWS Textract	OCR service — converts scanned images into text, structured data extraction from documents.
Amazon Comprehend	Extracts key phrases, entities, sentiment, and PII data from text.
AWS Intelligent Document Processing (IDP)	Automates the processing of unstructured documents (e.g., PDFs, invoices) using Textract, Comprehend, and A2I.
Amazon Lex	Service for building chatbots with natural conversational text or voice interfaces.
Amazon Transcribe	Speech-to-text service, including audio captioning.
Amazon Polly	Text-to-speech service for converting written text into spoken words.
Amazon Kendra	Intelligent document search engine with semantic understanding.
Amazon Personalize	Service for personalized product recommendations .
Amazon Translate	Language translation service.
Amazon Forecast	Provides time-series forecasting , e.g., inventory levels prediction.
Amazon Fraud Detection	Detects fraudulent activity, including online transactions and account takeovers.
Amazon Q Business	Generative AI-powered assistant for enterprise data processing and tasks.
Amazon Macie	PII data detection and anonymization service for data security.
SageMaker Clarify	Bias detection and explainability for machine learning models.
SageMaker Ground Truth	Provides human labeling for model training datasets.
Amazon Augmented AI (A2I)	Human review for low-confidence predictions during inference .
AWS Data Exchange	Access third-party datasets securely.
Amazon SageMaker JumpStart	Hub with pre-trained models and one-click deployments.
Fine-Tuning	Adjusting model weights using labeled data to improve task performance.
Domain adaptation fine-tuning	Tailor a model for a specific domain like legal or medical using small datasets.
Instruction-based fine-tuning	Fine-tuning a model to perform better on specific tasks , e.g., classification.
Continued-Pretraining	Using unlabelled data to expand a model's knowledge base .
Automatic Model Tuning (AMT)	Automatically tunes hyperparameters to improve model performance.
Data-Drift	Input data changes, but the relationship between inputs and outputs stays the same (e.g., new demographic).
Concept-Drift	The relationship between inputs and outputs changes, meaning the model's learned patterns no longer apply (e.g., new fraud patterns).
AWS Trusted Advisor	Provides recommendations for cost, performance, and security improvements.
Amazon Inspector	Automated security assessments for application workloads.

Term/Concept	Description
AWS PrivateLink	Secure private connections between VPCs and AWS services.
AWS Config	Monitors and records AWS configuration changes for compliance.
AWS CloudTrail	Logs and tracks AWS API calls for auditing.
BedRock GuardRails	Prevents inappropriate foundation model outputs and restricts risky content.
Postgres (Aurora or RDS)	SQL database with vector database support for similarity search.
Amazon DocumentDB	JSON store, MongoDB-compatible with vector database support.
Amazon Neptune	Graph database with vector search capabilities.
Amazon Neptune ML	Uses Graph Neural Networks (GNNs) to predict outcomes from graph data.
Amazon MemoryDB	In-memory database with fast vector search capabilities.
Amazon OpenSearch Service	Search service with vector database support for similarity search.
MSE (Mean Squared Error)	Average squared difference between predicted and actual values, lower MSE indicates better model performance.
RMSE (Root Mean Squared Error)	Square root of MSE, more interpretable; lower RMSE is better.
AWS Glue	ETL service for automating data extraction, transformation (e.g., Parquet), and loading into S3 for AI/ML workflows.
AWS Glue DataBrew	Visual tool for non-coders to clean and prep data via "recipes" with built-in quality checks for AI model training.
AWS Glue Transformations	Pre-built ETL functions (e.g., removing duplicates, filling missing values), supporting custom logic via Python/Spark for data prep.
SageMaker Canvas	No-code tool for data import, prep, and ML model building in a visual workflow, ideal for non-technical users.
SageMaker Data Wrangler	Unified interface for data cleaning, feature engineering, and prep for machine learning from multiple data sources.
AWS Glue Data Quality	Detects anomalies and suggests data quality rules to ensure clean, high-quality input for AI models.
AWS Glue Data Catalog	Centralized metadata store to manage and track schemas and ETL processes without storing actual data, for better data lineage management.

Confusion Matrix Evaluation Metrics:

Name	When to Use	Example
Precision	When minimizing false positives is crucial.	Spam detection (avoiding false positives): Avoiding flagging legitimate emails as spam.
Recall (TPR)	When minimizing false negatives is critical, focusing on identifying all positives .	Disease screenings (minimizing false negatives): Ensuring that no diseased patients are missed.
False Positive Rate (FPR)	When you need to avoid false alarms or incorrect positive predictions.	Security alarms (avoiding false positives): Avoiding alarms for non-threats or normal behavior.
Specificity (TNR)	When minimizing false positives and correctly identifying negatives is important.	Non-diseased patients (correct negatives): Correctly identifying healthy individuals in medical testing.
Accuracy	For general correctness across both positives and negatives .	Image classification (overall correctness): Correctly identifying both cats and dogs in a dataset.
F1 Score	When both precision and recall are important, and you want a balance between the two.	Document classification: Ensuring correct tagging of documents with minimal errors across both positives and negatives.